

1 The algorithm

```
type Image = PA (PA Int)
type Hist  = PA Int

hbalance :: Image -> Image
hbalance img =
  let a :: Hist
      a = joinD
        . mapD (\(as,a) -> mapS (plusInt a) as)
        . propagated plusInt 0
        . mapD (scanlS plusInt 0)
        . sparseToDenseD (plusInt gmax 1) 0
        . splitSparseD (plusInt gmax 1)
        . joinD
        . mapD tripletToATup2
        . segdSplitMerge 0
        . sortPS
        . concatPS
        $ img
      -- Histogram
      -- accu end

  n :: Int
  n = lengthPS a
      -- 2

  a0, divisor, gmax' :: Double
  a0 = int2Double . headPS $ a
      -- 3, variables for normalize and scale
  divisor = minusDouble (int2Double (lastPS a)) . int2Double . headPS $ a
  gmax' = int2Double $ gmax

  normScale :: Int -> Int
  normScale = floorDouble . (flip multDouble) gmax' . (flip divDouble) divisor . (flip minusDouble) a0 . int2Double
      -- 0, body of normalize and scale

  as :: Hist
  as = joinD . mapD (mapS normScale) . splitD $ a
      -- final mapping array
      -- 4, normalize and scale applied

  pixelReplicate :: Hist -> PA Hist
  pixelReplicate = concatPS . replPL (lengths (getSegd xs)) . replPS (lengthPS img)
      -- 0, artifact of NDP

in unconcatPS img
  . indexPL (pixelReplicate as)
  . concatPS
  $ img
      -- 5, apply. core of nested data parallelism here!
```

2 Work and Depth Table

- n sei die Anzahl der Bildpixel
- w sei die Bildbreite
- h sei die Bildhöhe
- p sei die Anzahl der PUs (gang members).

Table 1: Work and Depth complexities

function or variable	O(W)	O(D)
hbalance	$\max(n \log n, gmax)$	$\log \max(n, gmax)$
hist	$\max(n \log n, gmax)$	$\log n$
accu	$gmax$	$\log gmax$
as	$gmax$	1
pixelReplicate	1 (bzw. $n \cdot gmax$)	1
img'	n	1
sortPS	$n \log n$	$\log n$
concatPS	1	1
unconcatPS	1	1
headPS,lastPS	1	1
replPS $n \times$	n	1
replPL $ns \times xs$	$\text{sum}(ns) * \text{size}(xs)$	1
indexPL $is \times xs$	$\text{size}(xs)$	1
mapD tripletToATup2	$gmax \frac{p}{p}$	1
segdSplitMerge	$n \log n$	$\log n$
mapD fS xs	$W(fS, s) * \text{size}(xs)$	$W(fS, s) * \text{size}(xs) / p$
lengthPS	1	1
joinD xs	$\text{size}(xs)$	1 (communication overhead?)
splitD xs	$\text{size}(xs)$	1 (communication overhead?)
propagateD xs	1	$\log \text{size}(xs)$
sparseToDenseD	$gmax$	1
splitSparseD	$gmax$	1

3 Calculating "hbalance" entirely

Let $n = |img| = w \cdot h$

$$\begin{aligned}
 W(w \times h, gmax) &= W(hist) + W(accu) + W(as) + W(pixelReplicate) + W(img') \\
 &= O(\max(n \log n, gmax) + gmax + gmax + 1 + n) \\
 &\in O(\max(n \log n, gmax)) \\
 D(w \times h, gmax) &= \max\{hist, accu, as, pixelReplicate, img'\} \\
 &= \max\{\log n, \log gmax\} \\
 &\in O(\log \max(n, gmax))
 \end{aligned} \tag{1}$$

4 Calculating the histogram "hist"

Let $n = |img| = w \cdot h$

$$\begin{aligned}
W(w \times h) &= W(concatPS) + W(sortPS) + W(segSplitMerge) \\
&\quad + W(mapD, localSegdTup2) + W(joinD) \\
&= 1 + n \log n + n \log n + gmax + gmax \\
&= 1 + 2(n \log n + gmax) \\
&\in O(\max(n \log n, gmax)) \\
D(w \times h) &= \max\{concatPS, sortPS, ..., joinD\} \\
&\in O(\log n)
\end{aligned} \tag{2}$$

5 Calculating the accumulated histogram "accu"

$$\begin{aligned}
W(gmax) &= W(splitSparseD) + W(sparseToDenseD) + W(mapD, scanS) \\
&\quad + W(propagateD) + W(mapD, mapS, plusInt) + W(joinD) \\
&= gmax + gmax + gmax + 1 + gmax + gmax \\
&= 1 + 5 \cdot gmax \\
&\in O(gmax) \\
D(gmax) &= \max\{splitSparseD, sparseToDenseD, (mapD, scanS), ...\} \\
&\in O(\log gmax)
\end{aligned} \tag{3}$$

6 Calculating the mapping array "as"

$$\begin{aligned}
W(gmax) &= W(joinD) + W(mapD, normScale) + W(splitD) \\
&= gmax + gmax + gmax \\
&= 3 \cdot gmax \\
&\in O(gmax) \\
D(gmax) &= \max\{joinD, (mapD, normScale), splitD\} \\
&\in O(1)
\end{aligned} \tag{4}$$

7 Calculating by "pixelReplicate"

PixelReplicate is a function which executed strictly and literally would have a work complexity of $n \cdot gmax$. However, due to strong improvements in [WorkEfficient] this replication is simply cached locally such that the work complexity for this case is negligible and can be simplified to 1.

$$\begin{aligned}
W(gmax, w \times h) &\in O(1) \\
D(gmax, w \times h) &\in O(1)
\end{aligned} \tag{5}$$

8 Calculating the final balanced image "img'"

$$\begin{aligned} W(gmax, w \times h) &= W(unconcatPS) + W(indexPL) + W(concatPS) \\ &= 1 + n + 1 \\ &\in O(n) \\ D(gmax, w \times h) &= \max\{unconcatPS, indexPL, concatPS\} \\ &\in O(1) \end{aligned} \tag{6}$$

9 Other aspects e.g. sync-points, programmer workload, simplicity

- optimisations: removed sync-points(more local operations), tight distributed normalisation loop, distributed optimal prefix-sum, distributed groupP and sortP, nested data parallelism enables a depth lower than $O(\text{width})$
- progammer-workload: as much as for P_{nest}
- simplicity: many details, not indented for human readers, but shows optimisations clearly